

VanniKawachh: A Distributed AI Acoustic Intelligence and Autonomous Drone Response Network for Women Safety

A Thesis

Submitted in partial fulfilment of the requirements of the
Bachelor of Technology in Computer Science and Engineering

by

Group CSE_B_04

Shivansh Verma Saksham Sabadra Rudra Thakur Rohan Untawale

under the guidance of

Dr. Aditya Turankar

Department of Computer Science and Engineering
G. H. Rasoni College of Engineering, Nagpur

Session 2026-27

Certificate

This is to certify that the thesis entitled “*VanniKawachh: A Distributed AI Acoustic Intelligence and Autonomous Drone Response Network for Women Safety*”, submitted by Group CSE_B_04 (Shivansh Verma, Saksham Sabadra, Rudra Thakur, and Rohan Untawale) in partial fulfilment of the requirements of the Bachelor of Technology in Computer Science and Engineering, is a record of work carried out by the group under my guidance at the Department of Computer Science and Engineering, G. H. Rasoni College of Engineering, Nagpur, during the session 2026-27.

Dr. Aditya Turankar
Guide

Head of Department
Department of Computer Science and
Engineering

(Signatures, seal, and date to be completed.)

Declaration

We declare that this thesis is our own work, carried out by us as Group CSE_B_04 under the guidance of Dr. Aditya Turankar, and that all sources of material used (software libraries, protocol specifications, model checkpoints, regulatory texts, and surveyed literature) are cited in the References. The software described herein is published under the repository <https://github.com/SV-1411/drone.git>, and all experimental results reported are reproducible from that repository as described in Chapter 6. A large-language-model assistant was used during drafting for structure and consistency; all technical content derives from the implementation and its test logs, and every sentence was reviewed and accepted by the authors.

(Signatures and date to be completed.)

Abstract

Crimes against women often happen where help is hardest to reach: dark streets, forest stretches, campus outskirts, and parking areas. These places have few cameras, few patrols, and often no cellular coverage. Most safety technology aimed at this problem burdens the victim: panic-button apps, wearables, and carried devices all assume a charged, reachable device in the victim’s hand at the moment of attack. This thesis designs, implements, and validates **VanniKawachh** (“voice-shield”), a system that inverts the burden: the *infrastructure* listens, verifies, alerts, and responds, and the victim’s own voice is the only trigger she needs.

The system is a three-tier network. **Sensing nodes**, solar-powered pole-mounted ESP32-S3 boards with INMP441 I2S microphones, screen every audio frame on-device with a lightweight MFCC + CNN model (Stage 1, <50 ms per frame, recall-tuned), so no continuous audio ever leaves a pole. A **Raspberry Pi 5 hub** confirms each candidate event with PANNs deep audio tagging fused with PIR motion, ambient-light, and time-of-day evidence (Stage 2, precision-tuned). Confirmed alerts are AES-128-CTR encrypted, HMAC-authenticated, and replay-protected, and they carry the node’s surveyed coordinates from a registry rather than a live GPS fix. They travel over LoRa, with no SIM or cellular dependency, to a police dashboard, and simultaneously auto-dispatch a Pixhawk quadcopter that flies to the spot, records evidence during hover, descends to 3 m, drops a first-aid kit, and returns to launch. The response layer is the group’s previously built and SITL-verified autonomous flight stack (verified mode transition, a landing-interlocked mission queue, and a debounced severity-ordered failsafe arbiter), carried over unchanged.

Validation is staged, and we state clearly what has not been measured. The complete chain (synthesized distress audio → sealed packet → hub verification and fusion → dispatch → autonomous SITL flight with hover-record and kit release) runs end-to-end with zero hardware (Phase 0). The full-chain rehearsal passed with a fused severity of 0.88 and kit release commanded at 3.1 m, inside a 328 s mission, and the flight stack’s acceptance harness passes all 8/8 checks with 0.4 to 0.6 m terminal accuracy against a 5 m tolerance. Sixty-eight automated unit cases pin the safety logic, the packet cryptography, and the dispatch gating. Stage-1 model training on the microcontroller and hardware range/latency measurements are Phase 1 and 2 work in progress; the develop-

ment chain uses clearly labelled heuristic stand-ins, and this thesis reports no acoustic accuracy number it has not measured.

Keywords: women safety; acoustic event detection; edge AI; TinyML; PANNs; sensor fusion; LoRa; AES-128; autonomous UAV; verified dispatch; software-in-the-loop validation.

Contents

Certificate	i
Declaration	ii
Abstract	iii
1 Introduction	1
1.1 Motivation	1
1.2 The inverted-burden principle	2
1.3 Problem statement	2
1.4 Scope	3
1.5 Contributions	3
1.6 Thesis organization	3
2 Literature Review	5
2.1 Survey method	5
2.2 Bucket 1: victim-carried devices	5
2.3 Bucket 2: detection-only audio systems	6
2.4 The gap	6
2.5 Technology base	7
3 System Design	8
3.1 Three tiers	8
3.2 Design decisions and rationale	10
3.3 The mission lifecycle (response tier)	11
3.4 Data flow, end to end	11
4 Implementation: Sensing Node and Hub	13
4.1 Node firmware (<code>firmware/node/</code>)	13
4.2 Hub service (<code>hub/</code>)	15
4.3 Packet security (<code>hub/packets.py</code>)	16
4.4 Stage-2 verification (<code>hub/verifier.py</code>)	17

4.5	Evidence fusion (<code>hub/fusion.py</code>)	18
4.6	Dispatch (<code>hub/dispatcher.py</code>)	19
5	Implementation: The Response Layer (Safety-Verified Autonomous Flight)	20
5.1	Hazard analysis	21
5.2	Verified mode transition	21
5.3	The landing interlock	22
5.4	Failsafe arbitration	22
5.5	v2 extension: the DELIVERING phase	23
5.6	v2 extension: hover evidence recording	24
5.7	What the software refuses to own	24
6	Testing and Results	25
6.1	Methodology	25
6.2	Results: flight stack (v1 record, still current)	26
6.3	Results: Phase-0 full-chain rehearsal	26
6.4	Defect case studies (retained from v1)	27
6.5	What is not yet measured (Phase 1 and 2 work in progress)	28
7	Safety, Privacy, and Legal Compliance	29
7.1	Privacy by construction	29
7.2	Alert integrity: why spoofing is closed	30
7.3	Flight law (India, Drone Rules 2021)	30
7.4	Dispatch restraint	30
7.5	Radio-spectrum compliance	31
8	Conclusions and Future Work	32
8.1	Conclusions	32
8.2	Future work	32
A	LoRa Alert Packet Wire Format	37
B	Configuration Reference	38
C	Test Inventory	39

List of Figures

3.1	VanniKawachh three-tier system architecture: per-pole solar sensing nodes, the per-locality Raspberry Pi 5 hub, and the autonomous response drone.	8
3.2	Methodology flow: from the victim's shout through Stage-1 screening, Stage-2 verification and fusion, the sealed LoRa alert, to the police dashboard and the drone auto-dispatch.	12
4.1	The two-stage acoustic pipeline. Stage 1 on the ESP32-S3 (recall-tuned): 16 kHz audio $\rightarrow$ pre-emphasis and 32 ms framing $\rightarrow$ mel filterbank + log + DCT (13 MFCCs) $\rightarrow$ int8 CNN $\rightarrow$ alert + clip on a hit, on-device discard otherwise. Stage 2 on the hub runs PANNs tagging over the clip and fuses it with PIR/LDR/time before the dispatch gates.	13
4.2	The sealed 25-byte alert packet: cleartext header (magic, version, node_id, counter), AES-128-CTR ciphertext of the 8-byte payload, and the truncated HMAC-SHA256 tag.	16
4.3	Evidence fusion. The five evidence terms (Stage-2 audio score, Stage-1 confidence, PIR motion, darkness, and night) are combined with weights 0.60/0.15/0.10/0.08/0.07 into the severity S , and dispatch requires both $a \geq 0.50$ and $S \geq 0.60$ (Eqs. 4.17 and 4.18).	19
5.1	The mission executor's thirteen-state machine: IDLE $\rightarrow$ CONNECTING $\rightarrow$ WAITING_GPS $\rightarrow$ ARMING $\rightarrow$ TAKEOFF $\rightarrow$ ENROUTE $\rightarrow$ HOVERING $\rightarrow$ DELIVERING $\rightarrow$ RTL $\rightarrow$ LANDED $\rightarrow$ COMPLETED, with ABORTED and FAILED as abnormal terminals reachable from every flight phase.	20
5.2	The safety interlock chain: edge validation, queue admission, the pre-flight armed check, failsafe arbitration in every blocking loop, and the abort guarantee that blocks until disarm. These are the gates a dispatch must pass between network trigger and landing.	22

- 6.1 Altitude-time profile of the SITL response mission: takeoff to 15 m, en-route (896 m), hover + record, descent to the 3 m drop altitude with kit release commanded at 3.1 m, climb-out, return to launch, and landing. . 27

List of Tables

4.1	Hub service modules and their responsibilities.	15
5.1	Hazard analysis of the mission lifecycle.	21
6.1	SITL acceptance missions across the v1 development arc.	26
A.1	Sealed alert packet wire format (25 bytes).	37
B.1	Hub configuration (<code>hub/config.py</code>).	38
B.2	Flight-stack configuration (<code>flight_core/config.py</code> , v2 additions). . . .	38

Chapter 1

Introduction

1.1 Motivation

Consider where an assault actually happens. It is rarely on a well-lit arterial road under a CCTV camera. It is on the dark stretch between two streetlights, on the shortcut through a wooded campus edge, or in the far corner of a parking area at 23:40. These locations share three properties. First, they are *surveillance dead zones*: no camera, no patrol, often marginal or absent cellular coverage. Second, they are *response dead zones*: even a successful emergency call produces a ground unit minutes away through traffic. Third, and this is the property existing technology ignores, they are places where the victim is least able to operate a device: hands occupied, phone snatched, discharged, or simply out of reach in the seconds that matter.

The market's answer has been to instrument the victim. Panic-button mobile apps, smart wearables, GPS pendants, and IoT panic devices all place the trigger on the person at risk. The best of these report high accuracies, but the failure mode is structural: *a safety device the victim must carry, charge, and reach protects only the moments in which she can carry, charge, and reach it*. The research community's alternative, audio surveillance systems that detect screams, has demonstrated strong detection rates, but on server-scale compute, and it stops at detection. The output is a classifier score, not a located alert or a response on the ground.

What is missing is *infrastructure that owns the whole chain*: sensing that asks nothing of the victim, verification that suppresses false alarms well enough to act on, alert delivery that works precisely where cellular does not, and a first response that arrives in seconds rather than minutes. VanniKawachh is a design, an implementation, and a staged validation of that chain.

1.2 The inverted-burden principle

VanniKawachh’s single organizing idea is that the burden of triggering help must move from the victim to the environment. A scream, a cry, or a shouted “help” / “bachao” is a signal the victim produces anyway, under any circumstance, with no device. If pole-mounted infrastructure can hear it, verify it, and act on it, then the victim’s participation requirement drops to zero: no app, no wearable, no phone, no button.

The engineering consequences of that idea drive every design decision in this thesis:

- **Nodes must be cheap, solar, and everywhere.** So the node is an ESP32-S3 class microcontroller with a TinyML Stage-1 model rather than a streaming link to a server.
- **False alarms must be suppressed before anything expensive happens.** A drone launch is costly and attention-consuming, so a second, heavier verification stage runs at a hub, fused with environmental evidence.
- **The alert path must not depend on cellular coverage.** The target locations are dead zones by definition, so the alert travels over LoRa.
- **An alert that launches a drone must be unforgeable.** So every packet is encrypted, authenticated, and replay-protected.
- **The response aircraft must be trusted to fly unattended.** So the response layer is a flight stack whose safety properties are pinned by automated tests, built and validated before this concept existed (v1 of this project) and carried over unchanged.

1.3 Problem statement

Design and implement a distributed safety network in which pole-mounted acoustic nodes detect distress audio on-device, a locality hub confirms each event with deep audio analysis fused with environmental evidence, confirmed alerts travel encrypted over LoRa with the node’s surveyed coordinates to a police-facing dashboard, and an autonomous quadcopter is dispatched to the incident to record evidence and deliver a first-aid kit, such that (i) no continuous audio leaves any node; (ii) no alert can be spoofed or replayed into a drone launch; (iii) no dispatch occurs below explicit verification and severity thresholds; (iv) every flight-mode command is confirmed from autopilot telemetry, no mission can start against an airborne vehicle, and hazard responses obey debounce, severity, and escalation semantics; and (v) each of these properties is demonstrated by automated, reproducible tests before it is claimed.

1.4 Scope

In scope: the three-tier architecture; the node firmware structure and its Stage-1 pipeline; the complete hub service (packet cryptography, registry, Stage-2 verification, evidence fusion, dispatch gating); the response flight stack and its payload/camera extensions; end-to-end validation in software-in-the-loop simulation (Phase 0); and the specification of the hardware build phases. Out of scope for the prototype, and stated as such wherever relevant: the trained Stage-1 TFLite-Micro model (Phase-1 work in progress; the development chain uses a heuristic stand-in), hardware range and latency measurements (Phase 1 and 2), live video streaming to police, multi-node time-difference localization, and beyond-visual-line-of-sight flight operations.

1.5 Contributions

1. An **end-to-end architecture** (infrastructure sensing $\rightarrow$ two-stage verification $\rightarrow$ offline encrypted alerting $\rightarrow$ autonomous field response) where prior work covers at most one link of the chain (Chapters 2 and 3).
2. A **two-stage edge/hub acoustic pipeline**: a recall-tuned MFCC + CNN screen on the node (<50 ms budget) and a precision-tuned PANNs + sensor-fusion confirmation on the hub, with explicit thresholds gating dispatch (Chapter 4).
3. A **secure LoRa alert protocol**: a 25-byte sealed packet (AES-128-CTR + truncated HMAC-SHA256 + per-node monotonic replay counter, per-node keys derived from a master key) carrying an event class rather than coordinates, resolved against a surveyed node registry at the hub (Chapter 4, Appendix A).
4. A **safety-verified autonomous response layer** inherited from the group's v1 flight stack (verified mode transition, landing interlock, failsafe arbitration), extended with an evidence camera phase and a rule-bounded first-aid kit drop (descend to 3 m, release, fail $\rightarrow$ RTL) (Chapter 5).
5. A **staged validation methodology**: 68 automated unit cases, an 8/8 SITL acceptance flight, and a zero-hardware full-chain rehearsal (Phase 0) that proves the architecture before any soldering, with heuristic stand-ins labelled as such (Chapter 6).

1.6 Thesis organization

Chapter 2 surveys the literature and locates the gap. Chapter 3 presents the three-tier design and its rationale. Chapter 4 details the sensing node and hub implementation.

Chapter 5 details the response layer: the safety engineering of the flight stack, retained from v1 because it is unchanged and still load-bearing. Chapter 6 defines the validation methodology and reports results, including what has *not* yet been measured. Chapter 7 consolidates the safety, privacy, and legal posture. Chapter 8 concludes.

Chapter 2

Literature Review

2.1 Survey method

The group’s literature survey (twelve papers, 2023 to 2026, first catalogued in the Title Finalization Seminar record and reproduced with full bibliographic detail as references [1] to [12] of this thesis) covered the women-safety technology space across two communities: consumer/IoT safety devices and acoustic event detection. Every surveyed system falls into one of two buckets, and the buckets split on a single question: who carries the trigger?

2.2 Bucket 1: victim-carried devices

Mobile panic applications, smart wearables (bands, pendants, footwear sensors), and standalone IoT panic buttons [1, 5, 6, 9, 10, 11] place the sensing and triggering hardware on the victim. Pavithra et al. [1] present a discreet smartphone application with on-device audio recognition and emergency automation; Bharathi et al. [5] describe an IoT wearable combining GPS and accelerometer sensing and report **97.54%** detection accuracy with a 3.2 s response time and a 1.92% false-positive rate; Snehith et al. [6] and Hadkar et al. [9] present push-button/wearable IoT devices with locate-and-alert behaviour; Uganya et al. [10] describe a button-triggered GSM/GPS tracker; and Potturi et al. [11] automate the SOS trigger with GSM/GPS geo-location sharing. The stronger systems combine multiple modalities (accelerometer gestures, heart-rate anomalies, voice keywords), and GPS positioning with GSM/app-based alerting is near-universal in this bucket.

The limitation is availability rather than accuracy. Every system in this bucket fails the same way when the device is absent, damaged, discharged, snatched, or simply unreachable in the moment. These are exactly the conditions of a real assault. A 97.5%-accurate classifier on a device that is not in the victim’s hand protects nobody. The bucket also inherits cellular dependence: app and GSM alerting presume coverage that

the highest-risk locations often lack. Uganya et al.’s tracker [10], for example, does not work in the GSM dead zones where risk is highest.

2.3 Bucket 2: detection-only audio systems

The acoustic event detection literature [2, 3, 4, 7, 8] demonstrates that distress audio is machine-detectable at useful rates. Kim et al. [7] built an 11,921-sample large-scale scream dataset and found a CNN-Transformer to be the best of five evaluated models; the same group [2] later added a windowing CNN for scream temporal-interval prediction, improving F-measure by roughly three percentage points and equal-error rate by an order of magnitude. Sharma and Jebaseeli [3] report a scream-and-panic detector at **92%** accuracy with under 5% false positives; Fime et al. [8] achieve **95.51%** danger-sound accuracy with Noisereduce preprocessing and transfer-learned image backbones (InceptionV3, with MobileNetV2 as a lighter alternative) over mel-spectrogram inputs; and Srimathi et al. [4] show that scream detection can reduce reaction time by 40% and increase intervention rates by 30%. Ciaburro and Puyana-Romero’s systematic review of sound-event detection in smart cities [12] confirms both the maturity of detection methods and the near-absence of closed-loop response systems in the literature.

Three limitations recur across this bucket. First, **compute scale**: the reported models are server- or workstation-class [2, 7]; none runs on a solar pole budget. Second, **deployment binding**: evaluations are on curated datasets or tethered laboratory microphones, not distributed outdoor infrastructure. Third, and decisive for this thesis, **the chain stops at detection**: a positive classification is the *output* of these systems, with no location delivery, no alerting path engineered for dead zones, and no response on the ground [12].

2.4 The gap

Overlaying the buckets exposes the gap this project fills. Bucket 1 has response paths (alerts to guardians/police) but burdens the victim. Bucket 2 removes the victim’s burden but has no response path. **No surveyed system provides end-to-end infrastructure sensing → verified offline alerting → autonomous field response.** That chain, with each link individually feasible on published evidence, is VanniKawachh’s contribution, and each link imports a discipline from a different literature: TinyML keyword spotting for the node [14], large-scale pretrained audio tagging (PANNs [13]) for the hub, LPWAN practice for the alert path [16, 32], and the group’s own SITL-verified dispatch stack for the response [15, 22, 23, 24].

2.5 Technology base

TinyML on microcontrollers. TensorFlow Lite Micro [14] runs quantized CNNs in tens of kilobytes of RAM; the `micro_speech` class of MFCC-fronted keyword models is a proven template for sub-50 ms audio screening on ESP32-class silicon.

PANNs. Pretrained Audio Neural Networks [13] (CNN14 and lighter variants trained on AudioSet) provide calibrated per-class probabilities over hundreds of sound events, including the distress-relevant family (screaming, shouting, crying, yelling, wailing). Summing probability mass over that family yields a distress score without training a bespoke model, which suits a Stage-2 verifier on a Raspberry Pi 5.

LoRa. Semtech SX1278-class radios [32] deliver kilometre-scale links at roughly 1 to 5.5 kbps effective throughput with no network operator (the LoRaWAN specification [16] documents the underlying modulation and regional parameters). That is enough for a 25-byte alert and far too little for audio, which shapes the split transport design of §3.2.

The autopilot stack. ArduPilot and MAVLink [15] and software-in-the-loop simulation [24] are covered in the group’s v1 work; Chapter 5 summarizes what carries over. The v1 finding that motivates the whole project’s verification philosophy bears repeating: the standard client idiom for changing flight mode was observed to fail *silently* on ArduCopter 3.3 SITL: the library reports success while the autopilot ignores the command. The design answer, *a detection claimed is not a detection confirmed; a command sent is not a command adopted*, now governs every layer of VanniKawachh: sound, dispatch, and flight mode alike.

Regulation. India’s Drone Rules, 2021 (G.S.R. 589(E)) [19] govern the prototype’s flight operations (registration, zone constraints, visual line of sight); the Digital Personal Data Protection Act, 2023 [20] and plain privacy prudence govern the audio design; and the Wireless Planning and Coordination Wing’s delicensing framework [21] governs the radio links. Chapter 7 consolidates all three.

Chapter 3

System Design

3.1 Three tiers

Figure 3.1 shows the three-tier component architecture; the text schematic below annotates the same chain with the concrete hardware and software of each tier.




Figure 3.1: VanniKawachh three-tier system architecture: per-pole solar sensing nodes, the per-locality Raspberry Pi 5 hub, and the autonomous response drone.

```

+----- SENSING NODE (per pole, solar) -----+
| INMP441 I2S mic -> ESP32-S3: MFCC + tiny CNN (TFLM,<50ms) |
| PIR (HC-SR501) + LDR context ; Stage-1 hit -> alert+clip |
+-----+-----+-----+-----+-----+
| LoRa SX1278 (alert, AES-128)   ESP-NOW / WiFi (4 s clip) |
|                               |                             |
|                               |                             |
+-----+-----+-----+-----+-----+

```

```

+----- HUB (Raspberry Pi 5, per locality) -----+
| LoRa gateway (ESP32 + SX1278 on USB serial)      |
| Stage 2: PANNs (CNN14/CNN10) + PIR/LDR/time fusion score |
| Node registry: node_id -> surveyed (lat, lon)      |
| Police dashboard (live map + alarm) ; alert log    |
+-----+-----+
          POST /trigger {lat, lon, incident_type, priority}
          v
+----- RESPONSE DRONE (v1 stack, unchanged core) ----+
| trigger_api (FastAPI queue) -> mission_executor (13-state |
| FSM, verified mode setter, failsafe arbiter, landing      |
| interlock, obstacle keep-out routing)                  |
| v2: HOVERING + camera record -> DELIVERING (SG90 kit drop |
|      from 3 m; fail -> RTL) -> RTL                      |
+-----+-----+

```

The VanniKawachh chain in schematic form (cf. Figure 3.1). Each tier verifies before it acts: the node screens, the hub confirms, the flight stack confirms its own commands.

Tier 1: sensing node. One per pole. An ESP32-S3 continuously frames 16 kHz mono audio from an INMP441 I2S microphone, extracts MFCC features, and runs a tiny quantized CNN (TensorFlow Lite Micro, `micro_speech`-class) against the distress vocabulary: scream, cry, and the “help” / “bachao” keyword family. Budget: under 50 ms per frame. An HC-SR501 PIR sensor and an LDR provide motion and ambient-light context sampled alongside. A Stage-1 hit produces exactly two transmissions: a sealed 25-byte alert over LoRa, and a 4 s audio clip over ESP-NOW/WiFi. Otherwise the node transmits nothing.

Tier 2: hub. One Raspberry Pi 5 per locality. A gateway ESP32 with an SX1278 bridges LoRa to USB serial. The hub authenticates and decrypts each packet, resolves the node’s surveyed coordinates from a registry, waits briefly for the WiFi clip, re-scores the audio with PANNs, fuses the score with PIR/LDR/time evidence into a severity, and, only above two explicit thresholds, dispatches the drone and raises the police dashboard alarm. Every incident, dispatched or not, is logged.

Tier 3: response drone. The group’s v1 flight stack: a FastAPI [25] trigger surface, a priority mission queue, and a mission-executor state machine supervised by a failsafe arbiter. v2 adds a camera recording window during hover and a DELIVERING phase that descends to 3 m and releases a first-aid kit by servo. The stack’s safety core is untouched.

3.2 Design decisions and rationale

Fixed nodes carry no live GPS. Each pole is surveyed once at installation (NEO-6M or a phone fix); the hub’s registry maps `node_id` $\rightarrow$ (lat, lon, name). The LoRa packet therefore carries two bytes of identity instead of a coordinate pair. This is smaller, and it cannot be spoofed at the position level: a node has no GPS to jam, drift, or forge. An unknown `node_id` is dropped at the hub.

LoRa carries the alert; WiFi carries the audio. LoRa’s roughly 1 to 5.5 kbps cannot move a clip in useful time, so the transports split: the sealed alert goes over LoRa instantly (dead-zone-capable, kilometre-scale), and the 4 s verification clip follows over ESP-NOW/WiFi (~ 250 kbps, hundreds of metres line-of-sight).

The split is a direct consequence of LoRa’s physical layer. The duration of one LoRa symbol at spreading factor SF over bandwidth BW is

$$T_s = \frac{2^{\text{SF}}}{\text{BW}} \quad (3.1)$$

so higher spreading factors buy sensitivity at the price of airtime. The receiver’s minimum detectable signal is

$$P_{\min} [\text{dBm}] = -174 + 10 \log_{10}(\text{BW}) + \text{NF} + \text{SNR}_{\min} \quad (3.2)$$

where NF is the receiver noise figure and $\text{SNR}_{\min}$ the demodulation threshold for the chosen SF. Evaluating the standard LoRa time-on-air formula for the system’s 25-byte packet at SF9, 125 kHz bandwidth, coding rate 4:5 gives approximately **0.21 s** of airtime per alert; moving a 4 s, 16 kHz, 16-bit clip (128 kB) over the same link would take on the order of *minutes*. Hence alert-on-LoRa, clip-on-WiFi: the safety-critical 25 bytes ride the long-range dead-zone-capable channel, and the bulk audio rides the short-range fast one.

The design degrades gracefully: if the clip never arrives within the configured wait (8 s), the hub falls back to the Stage-1 confidence at a haircut ($\times 0.6$): enough to log always, and to dispatch only if the evidence is otherwise strong. Multi-node corroboration slots into the same fallback point as future work.

Stage 1 is recall-tuned; Stage 2 is precision-tuned. The node’s job is to *never miss* a real event; its false positives are expected and cheap, because the hub filters them with a model three orders of magnitude larger, plus environmental evidence. This division mirrors the economics: waking the hub costs milliwatts; launching the drone costs a mission. The philosophy is the v1 flight stack’s *verified dispatch* applied to sound.

Every LoRa packet is sealed. A spoofed packet would launch a drone. §4.3 details the AES-128-CTR + HMAC + replay-counter construction.

The flight core is untouched. All v1 safety machinery (verified mode setter,

failsafe arbiter, landing interlock, geofence, stall detection, obstacle keep-out routing) carries over unchanged. It is the reason the response half of this project already works, and Chapter 5 documents it as the response layer.

3.3 The mission lifecycle (response tier)

The executor’s state machine threads `IDLE` $\rightarrow$ `CONNECTING` $\rightarrow$ `WAITING_GPS` $\rightarrow$ `ARMING` $\rightarrow$ `TAKEOFF` $\rightarrow$ `ENROUTE` $\rightarrow$ `HOVERING` $\rightarrow$ **`DELIVERING`** $\rightarrow$ `RTL` $\rightarrow$ `LANDED` $\rightarrow$ `COMPLETED`, with `ABORTED` and `FAILED` as abnormal terminals: thirteen states in v2 (`DELIVERING` is the addition). Three v1 properties are preserved: no state waits for human input; every transition is timestamped into the mission log; every blocking loop polls both the failsafe arbiter and the operator-cancel flag, so abnormal termination is reachable from anywhere. The hover window doubles as the evidence-recording window; `DELIVERING` runs only when the mission requests a kit drop.

3.4 Data flow, end to end

Figure 3.2 traces one incident through the system, from the shout to the drone’s return; the numbered steps below give the corresponding software interfaces.

1. Node: Stage-1 hit $\rightarrow$ sealed alert (LoRa) + 4 s clip (WiFi).
2. Gateway ESP32: LoRa RX $\rightarrow$ one line per packet over USB serial. The gateway does no crypto and no parsing beyond framing; all intelligence stays on the Pi, where it can be updated without reflashing.
3. Hub pipeline: unseal (MAC, replay) $\rightarrow$ registry lookup $\rightarrow$ wait for clip $\rightarrow$ Stage-2 score $\rightarrow$ fusion $\rightarrow$ threshold gate $\rightarrow$ dispatch + log.
4. Drone stack: `POST /trigger {lat, lon, incident_type, priority}` $\rightarrow$ queue $\rightarrow$ mission with hover-record and kit drop $\rightarrow$ RTL.
5. Dashboard (React [26] + Leaflet [27] over OpenStreetMap [28] tiles): the incident appears on the police-facing map with severity, reasons, and mission id; the mission is trackable live.




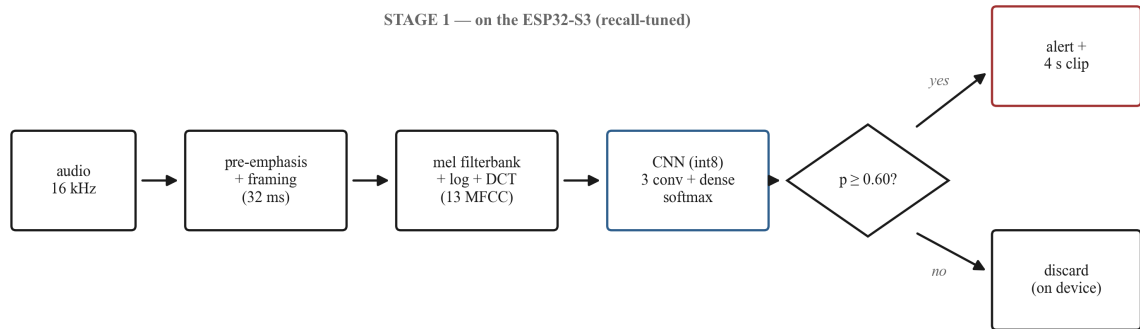
Figure 3.2: Methodology flow: from the victim’s shout through Stage-1 screening, Stage-2 verification and fusion, the sealed LoRa alert, to the police dashboard and the drone auto-dispatch.

Chapter 4

Implementation: Sensing Node and Hub

4.1 Node firmware (firmware/node/)

The sensing node is a single Arduino sketch for the ESP32-S3 [31] structured as the pipeline of §3.1: I2S capture from the INMP441 [33] at 16 kHz mono; framing and MFCC feature extraction; the TFLite-Micro Stage-1 model hook; PIR and LDR sampling; AES-sealed LoRa alert transmission (SX1278, LoRa library by Sandeep Mistry); and the ESP-NOW/WiFi clip upload to the hub’s clip server. The alert carries the event class (1 = scream, 2 = help_keyword, 3 = cry, 4 = crash), the Stage-1 confidence quantized to a byte, the PIR flag, the raw LDR level, the node battery percentage, and the monotonic packet counter. Figure 4.1 shows the Stage-1 pipeline as implemented on the node.



Stage 2 (hub): PANNs AudioSet tagging over the clip → fused with PIR/LDR/time → dispatch gates

Figure 4.1: The two-stage acoustic pipeline. Stage 1 on the ESP32-S3 (recall-tuned): 16 kHz audio → pre-emphasis and 32 ms framing → mel filterbank + log + DCT (13 MFCCs) → int8 CNN → alert + clip on a hit, on-device discard otherwise. Stage 2 on the hub runs PANNs tagging over the clip and fuses it with PIR/LDR/time before the dispatch gates.

The Stage-1 mathematics. The feature front-end is the classical MFCC chain. Each incoming frame is first pre-emphasized to flatten the spectral tilt of speech,

$$y[n] = x[n] - 0.97 x[n-1] \quad (4.1)$$

then windowed with a Hamming window of length N ,

$$w[n] = 0.54 - 0.46 \cos\left(\frac{2\pi n}{N-1}\right) \quad (4.2)$$

before the FFT. The power spectrum is pooled by a bank of M triangular filters spaced uniformly on the mel scale,

$$m = 2595 \log_{10}\left(1 + \frac{f}{700}\right) \quad (4.3)$$

and the log energy of each filter output is taken,

$$e_j = \log\left(\sum_k H_j(k) |X(k)|^2\right), \quad j = 1, \dots, M \quad (4.4)$$

where H_j is the j -th triangular filter. A DCT-II decorrelates the log energies into the cepstral coefficients

$$c_i = \sum_{j=1}^M e_j \cos\left[i\left(j - \frac{1}{2}\right) \frac{\pi}{M}\right], \quad i = 0, \dots, 12 \quad (4.5)$$

of which the first **13** are retained as the frame’s feature vector.

The classifier head is a small CNN ending in a softmax over the K event classes,

$$p_k = \frac{e^{z_k}}{\sum_{j=1}^K e^{z_j}} \quad (4.6)$$

trained (Phase 1) by minimizing the cross-entropy loss

$$\mathcal{L} = - \sum_{k=1}^K y_k \log p_k \quad (4.7)$$

against one-hot labels y . For deployment under TensorFlow Lite Micro the weights and activations are quantized to int8 with the affine map

$$x_q = \text{round}\left(\frac{x}{s}\right) + z \quad (4.8)$$

where s is the per-tensor scale and z the zero point. This step shrinks the model into the

ESP32-S3’s RAM budget and the <50 ms frame deadline.

Status honesty. The firmware’s capture, sensing, sealing, and transport paths are implemented; the Stage-1 *model* is a hook. Training the quantized MFCC + CNN on scream/cry/keyword data and flashing it is Phase-1 work in progress (§6.5). Nothing in this thesis claims an on-device detection accuracy.

4.2 Hub service (hub/)

The hub is a Python package on the Raspberry Pi 5 [30], one module per responsibility (Table 4.1).

Table 4.1: Hub service modules and their responsibilities.

Module	Responsibility
hub/config.py	Env-driven settings: serial port, master key, thresholds, drone API URL
hub/node_registry.py	<code>node_id</code> $\rightarrow$ (lat, lon, meta), JSON-backed; unknown id $\Rightarrow$ packet dropped; tracks each node’s last counter
hub/packets.py	25-byte packet format; AES-128-CTR seal/unseal, truncated HMAC-SHA256, replay counter (§4.3)
hub/lora_gateway.py	Reads the gateway ESP32’s USB serial stream; <code>--sim</code> mode substitutes synthetic packets
hub/clip_server.py	Receives the nodes’ 4 s WAV clips over WiFi
hub/verifier.py	Stage-2 scoring: PANNs backend, or an energy-heuristic dev fallback (§4.4)
hub/fusion.py	Severity fusion of audio score with PIR/LDR/time evidence (§4.5)
hub/pipeline.py	The gate: alert $\rightarrow$ verify $\rightarrow$ fuse $\rightarrow$ dispatch decision; no dispatch below threshold
hub/dispatcher.py	POSTs <code>/trigger</code> to the drone stack with the node’s surveyed coordinates
hub/main.py	Entrypoint: <code>python -m hub.main (serial) or --sim</code>

The pipeline (`process_packet`) executes the full chain for one sealed packet: authenticate + decrypt + replay-check; registry lookup (bump the node’s counter only after acceptance); wait up to 8 s for the clip at `hub/clips/<node_id>_<counter>.wav`; Stage-2 score the clip, or, if the clip never arrives, substitute the degraded audio score

$$a = 0.6 c \tag{4.9}$$

where c is the Stage-1 confidence (the no-clip haircut of §3.2); fuse; then gate: dispatch requires **both** `audio_score` ≥ 0.50 (`VERIFY_THRESHOLD`) **and** `severity` ≥ 0.60 (`DISPATCH_THRESHOLD`). Below either threshold the incident is logged with its reasons and no drone flies. Every incident, dispatched or not, is appended to the incident list that feeds the dashboard.

4.3 Packet security (hub/packets.py)

The threat model is simple: a spoofed packet launches a drone; an eavesdropped packet reveals an incident in progress; a replayed packet re-launches a drone at an attacker's chosen time. The wire format (25 bytes, which fits in one LoRa frame; Figure 4.2, full layout in Appendix A) answers all three:

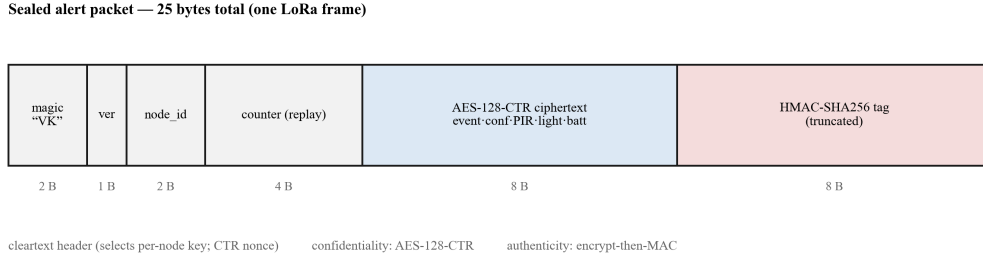


Figure 4.2: The sealed 25-byte alert packet: cleartext header (magic, version, node_id, counter), AES-128-CTR ciphertext of the 8-byte payload, and the truncated HMAC-SHA256 tag.

- **Confidentiality:** the 8-byte payload (event, confidence, PIR, light, battery) is AES-128-CTR encrypted. The CTR nonce is derived from the cleartext header (magic, version, node_id, counter), unique per packet as long as the counter is monotonic.
- **Authenticity:** an 8-byte MAC (HMAC-SHA256 over header + ciphertext, truncated) using the per-node key. A bad MAC is dropped before decryption is trusted.
- **Replay protection:** the uint32 counter is monotonic per node; the hub rejects any counter $\leq$ the last accepted value for that node.
- **Key management:** each node key is derived as $\text{HMAC-SHA256}(\text{master_key}, \text{"node:"} + \text{id})[0:16]$, so provisioning a node requires only the master key and its id; the hub holds one secret.

`node_id` and `counter` travel in cleartext by necessity (the id selects the key; the counter builds the nonce). Neither is sensitive, and both are covered by the MAC.

Formally, the construction is as follows. Node n 's key is derived from the hub's master key K_m as

$$K_n = \text{HMAC-SHA256}(K_m, \text{"node:"} + n)[0:16] \quad (4.10)$$

so the hub holds one secret and each node holds only its own derived key (HMAC per RFC 2104 [18]). The 8-byte payload P is encrypted in counter mode,

$$C = P \oplus E_{K_n}(\text{IV}) \quad (4.11)$$

where E is AES-128 [17] and the initial counter block IV is built from the cleartext header (magic, version, node_id, counter). The IV is unique per packet as long as the node's counter is monotonic. The authentication tag is an encrypt-then-MAC over everything on the wire,

$$\text{tag} = \text{HMAC-SHA256}(K_n, \text{header} \parallel C) [0:8] \quad (4.12)$$

and the hub accepts a packet if and only if

$$\text{tag valid} \wedge \text{counter} > \text{counter}_{\text{last}} \quad (4.13)$$

for that node, bumping $\text{counter}_{\text{last}}$ only after acceptance. The truncated 8-byte tag leaves a blind-forgery success probability of 2^{-64} per attempt, which is negligible at LoRa alert rates; the trade-off would be revisited if the channel were ever widened.

4.4 Stage-2 verification (hub/verifier.py)

Two interchangeable backends behind one interface (`verify_wav(path) → score in [0, 1]`):

- **PANNs (production).** The pretrained AudioSet tagging model (`panns-inference`, CNN14 by default; CNN10 or a MobileNet variant if the Pi is slow). The distress score is the summed probability over the distress-relevant AudioSet classes (screaming, shouting, yelling, crying, wailing, groaning, whimpering), clamped to 1.0. No bespoke training is required: the verifier leans on AudioSet scale, which a per-locality hub can afford to run and a pole cannot.
- **Energy heuristic (dev/SITL fallback).** Loud + high-spectral-centroid + bursty audio scores high (weights 0.45/0.35/0.20). This backend exists so the entire chain runs on any machine with no torch installation. It is **not** a claim of detection accuracy, and every result produced with it is labelled as fallback, including the Phase-0 numbers of Chapter 6.

The fallback scorer is worth stating exactly, because it gates the Phase-0 rehearsal. Over the clip's samples $x[n]$, $n = 1, \dots, N$, it computes the loudness as the root-mean-square level

$$\text{RMS} = \sqrt{\frac{1}{N} \sum_{n=1}^N x[n]^2} \quad (4.14)$$

and the spectral centroid (the “brightness” of the clip, high for screams)

$$C = \frac{\sum_k f_k |X(k)|}{\sum_k |X(k)|} \quad (4.15)$$

where $X(k)$ is the magnitude spectrum at frequency f_k . Together with a burstiness term (burst $\in [0, 1]$, from the envelope’s peak-to-mean structure) the score is

$$\text{score} = 0.45 \min\left(1, \frac{\text{RMS}}{0.15}\right) + 0.35 \text{clip}\left(\frac{C - 400}{1600}, 0, 1\right) + 0.20 \text{burst} \quad (4.16)$$

i.e. full loudness credit at RMS ≥ 0.15 full scale, and centroid credit ramping linearly from 400 Hz to 2 kHz.

4.5 Evidence fusion (`hub/fusion.py`)

A night-time scream in a dark spot with motion nearby carries more weight than a daytime shout on a busy road. With a the Stage-2 audio score, c the Stage-1 confidence, $p \in \{0, 1\}$ the PIR motion flag, $L \in [0, 255]$ the LDR light level, and $n \in \{0, 1\}$ the night indicator (1 between 20:00 and 06:00), the fused severity is the weighted sum (Figure 4.3)

$$S = 0.60 a + 0.15 c + 0.10 p + 0.08 d + 0.07 n, \quad d = 1 - \frac{L}{255} \quad (4.17)$$

clamped to $[0, 1]$. Audio dominates by design; the environmental terms nudge. The dispatch gate of §4.2 is then the conjunction

$$\text{dispatch} \iff a \geq 0.50 \wedge S \geq 0.60 \quad (4.18)$$

and the mission priority is

$$\text{priority} = \text{high} \iff S \geq 0.75 \vee (a \geq 0.6 \wedge p = 1) \quad (4.19)$$

otherwise `normal`.

Every fusion emits a human-readable reasons string (`audio=... stage1=... pir=... dark=... night=...`) that travels to the log and dashboard; an operator can always see *why* a drone launched or an incident was held. The weights are prototype values, stated as such, to be tuned against Phase-1 bench data.




Figure 4.3: Evidence fusion. The five evidence terms (Stage-2 audio score, Stage-1 confidence, PIR motion, darkness, and night) are combined with weights 0.60/0.15/0.10/0.08/0.07 into the severity S , and dispatch requires both $a \geq 0.50$ and $S \geq 0.60$ (Eqs. 4.17 and 4.18).

4.6 Dispatch (`hub/dispatcher.py`)

On a gate-passing incident the dispatcher POSTs the v1 trigger API: `target` = the node's surveyed coordinates, `incident_type` from the event class, `priority` from fusion, `deliver_kit` set. The drone stack's own edge validation (geofence, altitude bounds, queue admission) applies unchanged. The hub is a client of the same hardened API any operator would use, not a privileged backdoor.

Chapter 5

Implementation: The Response Layer (Safety-Verified Autonomous Flight)

The response layer is the group’s v1 flight stack, built and validated before the VanniKawachh pivot, and retained unchanged because its properties are exactly what an unattended women-safety responder needs. This chapter preserves the v1 technical record, which is still accurate, and adds the two v2 extensions (camera, payload). Figure 5.1 shows the executor’s thirteen-state mission state machine (§3.3), including the v2 DELIVERING phase.

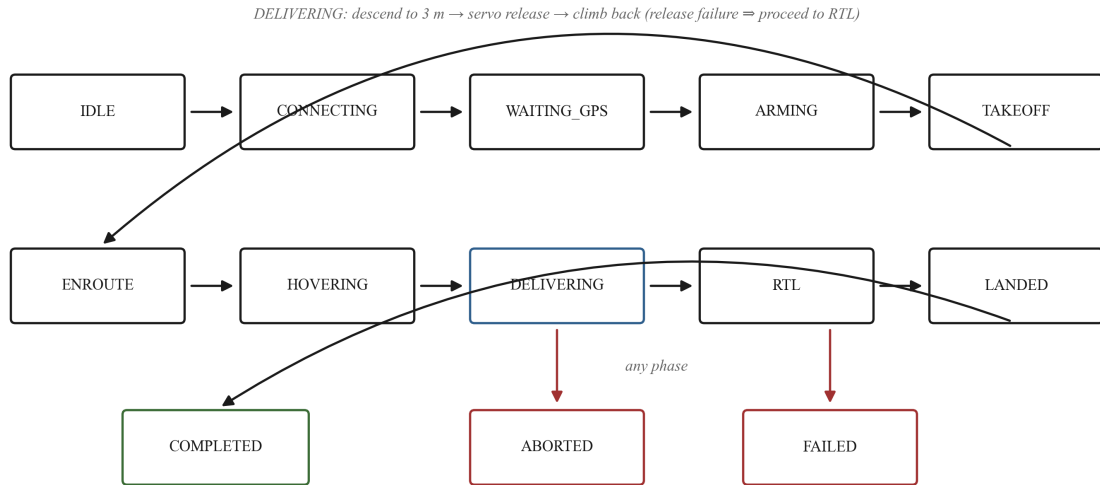


Figure 5.1: The mission executor’s thirteen-state machine: IDLE → CONNECTING → WAITING_GPS → ARMING → TAKEOFF → ENROUTE → HOVERING → DELIVERING → RTL → LANDED → COMPLETED, with ABORTED and FAILED as abnormal terminals reachable from every flight phase.

5.1 Hazard analysis

A compact HAZOP-style pass over the mission lifecycle yields the hazard set the design must close (Table 5.1).

Table 5.1: Hazard analysis of the mission lifecycle.

ID	Hazard	Worst credible outcome	out- Closed by
H1	Mode command silently rejected	Stranded/uncontrolled aircraft; unrecalable mission	§5.2 verified setter
H2	Emergency command (RTL/LAND) not delivered	Failsafe ineffective at the moment it matters	§5.2 + cross-action fallback
H3	New mission starts while airborne	Takeoff commanded to flying vehicle; undefined behaviour	§5.3 interlock
H4	Transient GPS dropout treated as loss	Unnecessary landing on unsafe ground	§5.4 debounce
H5	Critical battery during RTL	Aircraft presses home and falls short	§5.4 mid-return escalation
H6	Demand downgraded (LAND→RTL)	Wrong recovery flown	§5.4 monotone severity
H7	Target beyond geofence accepted	Predictable mid-air abort; wasted battery at distance	edge validation
H8	Wind stall / rejected goto	Battery exhausted loitering	leg-stall detector
H9	Process shutdown mid-flight	Aircraft abandoned under autopilot defaults	shutdown RTL
H10	Dispatch by unauthorized party	Aircraft weaponized by anyone on the network	token auth + sealed LoRa path (Ch. 4)
H11	Kit-release failure at low altitude	Aircraft loiters at 3 m over an incident scene	§5.5 fail → RTL rule


5.2 Verified mode transition

Every flight-mode command, nominal or emergency, routes through one routine, `_set_mode_confirmed`: issue through the high-level interface; loop until deadline reading the autopilot’s HEARTBEAT-derived mode; every 700 ms of non-confirmation re-issue through `COMMAND_LONG(MAV_CMD_DO_SET_MODE)` and the legacy `SET_MODE` encoding and re-poke the high-level setter; success iff the *autopilot reports* the requested mode. Failure on the abort path triggers the cross-action fallback: an unconfirmable RTL becomes a LAND attempt and vice versa, because some confirmed recovery beats an optimal unconfirmed one. Idempotence of mode-setting makes the retry loop safe; reading confirmation

from the autopilot’s own report makes it sound. The need is empirical: the standard client idiom for entering GUIDED is silently ignored by the ArduCopter 3.3 simulator while reporting success (H1). This was observed in v1’s first flight, and it is the origin of the whole project’s verification philosophy. The bare setter appears nowhere outside the routine; a reviewer can verify the property by grepping for mode assignments.

5.3 The landing interlock

Two mechanisms, redundant by intent (Figure 5.2). The **abort guarantee**: every abnormal termination (failsafe, recall, exception with an airborne vehicle) commands its action through §5.2 and then blocks, polling armed state and relative altitude, until disarm or a 240 s bound, before the queue regains control. The **pre-flight check**: a dequeued mission re-reads the armed flag and waits (bounded, 120 s) or fails rather than arm an armed vehicle. Together they make the invariant (*the queue can never start a flight against an airborne vehicle*) guaranteed by the design, not left to chance. The same property is defended at process boundaries: shutdown with an armed vehicle commands RTL (verified) before the MAVLink link drops (H9).



Invariant: no mission is ever commenced against an armed or airborne vehicle.

Figure 5.2: The safety interlock chain: edge validation, queue admission, the pre-flight armed check, failsafe arbitration in every blocking loop, and the abort guarantee that blocks until disarm. These are the gates a dispatch must pass between network trigger and landing.

5.4 Failsafe arbitration

The arbiter evaluates four condition families at 1 Hz (battery low (20%) and critical (10%), GPS fix validity, geofence distance, mission wall-clock) and maintains exactly one demanded action under three rules. **Monotone severity**: demands only escalate (NONE < RTL < LAND); H6 closed by construction. **Debounce**: GPS loss fires only after N (default 3) consecutive bad samples, reset on recovery; it demands LAND because

a return path is non-navigable without positioning; H4 closed with detection latency bounded at N seconds. **Fire-once:** each named hazard emits once per mission, with re-emission permitted only as escalation, so the incident log is not flooded with repeats. The executor couples to the arbiter inside every blocking loop, *including the RTL loop*, which re-reads the demand each second and swaps to LAND on escalation (H5).

Two pieces of mathematics underpin the arbiter’s inputs. The geofence evaluator measures the vehicle’s great-circle distance from home with the haversine formula: for latitudes φ_1, φ_2 and longitude difference $\Delta\lambda$,

$$d = 2R \arcsin \sqrt{\sin^2 \frac{\Delta\varphi}{2} + \cos \varphi_1 \cos \varphi_2 \sin^2 \frac{\Delta\lambda}{2}}, \quad R = 6371 \text{ km} \quad (5.1)$$

This is the same computation used by edge validation (H7) and by the acceptance harness’s closest-approach metric. The GPS-fix validity the arbiter debounces is itself the output of the autopilot’s extended Kalman filter, which maintains the state estimate $\hat{x}$ through the standard predict/update pair

$$\hat{x}_{k|k-1} = F_k \hat{x}_{k-1|k-1}, \quad P_{k|k-1} = F_k P_{k-1|k-1} F_k^\top + Q_k \quad (5.2)$$

$$\hat{x}_{k|k} = \hat{x}_{k|k-1} + K_k (z_k - H_k \hat{x}_{k|k-1}), \quad K_k = P_{k|k-1} H_k^\top (H_k P_{k|k-1} H_k^\top + R_k)^{-1} \quad (5.3)$$

fusing GPS, IMU, and barometer. The companion computer deliberately does not re-implement this estimator; it consumes the autopilot’s fix type and debounces it, trusting the EKF for fusion and itself for policy.

5.5 v2 extension: the DELIVERING phase

(`flight_core/payload_release.py`)

After the hover window, a mission flagged `deliver_kit` enters DELIVERING: the executor commands a descent over the incident point to the configured drop altitude (default **3.0 m**, tolerance 0.7 m, bounded at 45 s; on timeout it drops from current altitude rather than loiter), releases the first-aid kit, then climbs back to cruise altitude for a clean RTL. The release itself is an SG90 servo on a Pixhawk [29] AUX output (servo channel 9), commanded via `MAV_CMD_DO_SET_SERVO`, a protocol-level command with no client-library dependency, so it behaves identically over SITL (where the simulator accepts and logs it) and hardware (where the hook physically opens). Open PWM 1900, hold for a 2 s settle, re-close at PWM 1100.

The 3 m drop altitude comes from simple ballistics. Released from rest at height h ,

the kit falls for

$$t = \sqrt{\frac{2h}{g}} \approx 0.78 \text{ s} \quad (h = 3 \text{ m}) \quad (5.4)$$

and a horizontal wind of speed v displaces it by at most

$$\Delta x \approx v t \approx 1.6 \text{ m} \quad (v = 2 \text{ m/s}) \quad (5.5)$$

So from 3 m, even in a 2 m/s breeze, the kit lands within arm's reach of the incident point, and the impact energy of a soft-packed kit is small. From transit altitude (15 m) the fall takes 1.75 s and the same wind drifts it ~ 3.5 m onto uncertain ground, which is why nothing is ever dropped from transit.

The governing rule (H11): **a failed release is never a reason to loiter.** The failure is logged and reported, and the drone proceeds to RTL regardless. Physical drop confirmation (a payload microswitch) is future work; the current return value confirms command acceptance. The DELIVERING loop polls the failsafe arbiter and abort flag like every other phase, so a battery or GPS demand pre-empts the drop.

5.6 v2 extension: hover evidence recording (`flight_core/camera_recorder.py`)

The recorder starts when the mission enters HOVERING and stops after DELIVERING, writing an mp4 tagged with the mission id into the mission log directory. This file is the evidence artifact for the incident record. On machines without a camera (SITL, dev laptops) it is a no-op stub, so the mission flow is byte-identical either way; on the airframe it drives the Pi Camera Module 3. Recording is mission-scoped by construction: the camera runs only over the incident, never in transit, and live streaming to police is explicitly future work (Chapter 8).

5.7 What the software refuses to own

Two authorities stay outside the stack by design. The hardware RC link (pilot mode-switch and kill) overrides anything this software commands; the last line of defence is hardware, not software. And ArduPilot's own pre-arm checks and firmware failsafes stay at stock values on hardware: the SITL-only relaxation is dead code unless `SITL_MODE=1`, and the documentation treats setting it on a real aircraft as a defect, not an option.

Chapter 6

Testing and Results

6.1 Methodology

Validation follows the project’s build phases (Phase 0 $\rightarrow$ 4), with one governing rule inherited from the project plan: *every safety claim in this document is pinned by an automated test before it is written down, and every number produced by a stand-in is labelled as such*. Three tiers exist today:

Tier 1: unit suite (68 cases, no simulator, seconds to run). Three files. `tests/test_units.py` (47 collected cases) drives the flight stack’s safety logic with a synthetic vehicle: every arbiter rule of §5.4 (low $\rightarrow$ RTL; critical $\rightarrow$ LAND including over an in-progress RTL; never downgrade; $N - 1$ bad GPS samples $\rightarrow$ no trigger, $N \rightarrow$ trigger, recovery resets; fire-once; geofence; timeout), queue semantics (priority order, depth rejection, cancel of queued and running missions, history pruning), persistence round-trip with crash-orphan marking, every edge-validation rejection bound, and configuration semantics. `tests/test_hub.py` (14 cases) pins the v2 chain: packet seal/unseal round-trip, MAC rejection on tamper, replay rejection on stale counters, registry lookup and unknown-node drop, fusion scoring and priority mapping, pipeline gating (**no dispatch below threshold** is asserted, not assumed), and the dispatcher’s payload shape. `tests/test_obstacle_avoidance.py` (7 cases) covers keep-out routing geometry. The suites discriminate: run against the pre-hardening v1 implementation, the debounce, cancel, and interlock cases fail. The tests encode the safety claims rather than mirroring the code.

Tier 2: SITL acceptance flight. The harness boots `dronekit-sitl` `copter-3.3` and the API as child processes, dispatches the 896 m test mission (altitude 15 m, dwell 5 s), polls telemetry at 1 Hz, and asserts eight properties: SITL listening, API listening, vehicle connected, armed, took off ($\geq 80\%$ of target altitude), reached target (closest approach ≤ 5 m), returned home (≤ 10 m, required), landed. No tolerance widening is applied.

Tier 3: Phase-0 full-chain rehearsal (scripts/demo_phase0.py). The complete VanniKawachh chain with zero hardware: a synthesized distress WAV (loud, high-pitched, bursty; a stand-in for Phase-1 recorded test audio) stands in for the microphone; a simulated node 600 m from home seals a real 25-byte alert packet with the real cryptography; the real hub pipeline unseals it, waits for the clip, scores it with the energy-heuristic fallback backend, fuses, gates, and dispatches; the real drone stack flies the SITL mission with the hover-record window and the DELIVERING kit drop. Everything except the audio source, the Stage-2 backend, and the physics is production code.

6.2 Results: flight stack (v1 record, still current)

Six end-to-end SITL missions span the v1 development arc (Windows 11, Python 3.11.9); Table 6.1 summarizes the acceptance metrics.

Table 6.1: SITL acceptance missions across the v1 development arc.

Metric	Run 2	Run 3	Run 4	Run 5	Run 6 (hardened)
Wall-clock (s)	503.7	321.6	321.3	330.7	331.3
Closest approach (m)	0.4	0.5	0.4	0.4	0.6
Final distance from home (m)	0.1	0.0	0.0	0.0	0.0 to 0.2
Verdict	PASS [†]	PASS	PASS	PASS	PASS (8/8)

[†] harness predicate bug (§6.4, case 2); verdict unaffected.

Terminal accuracy of 0.4 to 0.6 m sits an order of magnitude inside the 5 m tolerance. It is bounded by the autopilot’s loiter behaviour, not the dispatch layer. Wall-clock variance across runs 3 to 6 is <3%, dominated by simulator boot.

6.3 Results: Phase-0 full-chain rehearsal

The rehearsal **passed** end to end. Chain trace, as logged:

- Sealed alert from simulated node (event **scream**, PIR active, dark LDR reading) authenticated, decrypted, and replay-checked by the real packet code; registry resolved the surveyed coordinates.
- Stage-2 fallback backend scored the synthesized clip above the 0.50 verification threshold.
- **Fused severity: 0.88** → priority **high**, clearing the 0.60 dispatch threshold.
- Dispatch POSTed; the SITL mission ran the full thirteen-state lifecycle: takeoff, obstacle-aware transit, hover with the recording window active (no-op recorder

in SITL), then DELIVERING (descent commanded to 3.0 m, **kit release commanded at 3.1 m** relative altitude, hook re-closed), then climb-out, RTL, land, disarm.

- Mission wall clock: **328 s**; closest approach 0.4 m; acceptance properties 8/8.

Figure 6.1 plots the mission’s altitude-time profile: the 15 m cruise out over the 896 m leg, the hover-record window, the descent spike to the 3 m drop line for the kit release, and the return.




Figure 6.1: Altitude-time profile of the SITL response mission: takeoff to 15 m, enroute (896 m), hover + record, descent to the 3 m drop altitude with kit release commanded at 3.1 m, climb-out, return to launch, and landing.

Unit tier: **68/68 pass** on the current implementation.

This proves two things: the *architecture* is sound and the *integration* is real. Every interface a hardware phase will use (packet bytes, clip convention, thresholds, trigger payload, servo command) was exercised in its production form. It deliberately does not prove acoustic detection performance (see §6.5).

6.4 Defect case studies (retained from v1)

Three defects found during v1 are retained as evidence for the design philosophy, because they now protect the whole VanniKawachh chain.

Case 1: the silent mode rejection (run 1). State machine reached TAKEOFF, altitude stayed 0.0 m, vehicle disarmed itself. Log forensics showed mode still STABILIZE despite the client reporting GUIDED. Root cause: unverified setter (H1). Fix: §5.2. Lesson: *the autopilot’s report is the only truth; the client’s state is a wish*. The two-stage acoustic pipeline applies the same lesson to Stage-1 claims.

Case 2: the racing test predicate (run 2). The harness inferred “connected” from a telemetry state string and raced a transition. Fix: an explicit boolean on `/health`. Lesson: tests must consume *interfaces*, not coincidences of internal state.

Case 3: the unguarded abort path (found by review, closed before it fired). The original abort used the bare setter (the exact call proven unreliable in Case 1) and returned with the vehicle airborne (H2+H3 compound). The unit tier now contains the regression tests that would have caught both. Lesson: *the emergency path must be held to a higher standard than the nominal path, and it is the one most easily left to a lower one.*

6.5 What is not yet measured (Phase 1 and 2 work in progress)

Stated plainly, because the credibility of the measured numbers depends on the honesty of the unmeasured ones:

- **No Stage-1 accuracy number exists.** The TFLite-Micro model is a hook; training it on scream/cry/keyword data (“help”, “bachao”) and flashing it to the ESP32-S3 is Phase 1. The <50 ms figure is the design budget for a `micro_speech`-class model, not a measurement.
- **No Stage-2 accuracy number exists for this deployment.** The Phase-0 score came from the labelled energy-heuristic *fallback*, scoring a *synthesized* clip. PANNs’ published AudioSet performance [13] motivates the backend choice but is not a claim about this system’s field precision. Phase-1 bench work measures Stage-2 latency and end-to-end false-positive rate on real street noise, and tunes the fusion weights.
- **No radio range or loss figures exist.** LoRa range (urban/open) and packet loss vs. spreading factor, and ESP-NOW clip-delivery reliability, are Phase-2 measurements.
- **No hardware flight has occurred.** All flight results are SITL (ArduCopter 3.3, a 2015 release; its command-delivery faults forced the defensive design, but it remains an old simulator); Phase 3’s staged progression (bench → props-off → manual hover → guided leg → full auto, VLOS) governs the transition.

These are the numbers the Phase 1 and 2 bench campaigns exist to produce, and the papers derived from this thesis will carry them only once measured.

Chapter 7

Safety, Privacy, and Legal Compliance

7.1 Privacy by construction

No continuous recording or transmission exists anywhere in the system, structurally: audio is processed in-place on the node, frame by frame, and a frame that does not trip Stage 1 is discarded: nothing stored, nothing transmitted. Only event-triggered clips of at most 5 s ever leave a node, and the accompanying alert is encrypted. The node cannot stream even if compromised at the configuration level, because the transport budget (LoRa) cannot carry audio and the clip path is event-gated in firmware. On the drone, the camera records only over the incident scene during the hover/delivery window, mission-tagged, for evidentiary use, never in transit. This posture is stated verbatim in the project plan and repeated on every public-facing description of the system, because a listening-infrastructure project earns deployment consent with exactly this property.

The same posture is what India’s **Digital Personal Data Protection Act, 2023** [20] asks of a deployment. The Act’s data-minimisation and purpose-limitation principles are satisfied structurally rather than by policy: on-device processing means no personal data is collected at all for the overwhelming majority of frames (a discarded frame is never “processed” off the node); the only audio that ever constitutes stored data is the event-triggered clip of at most 5 s, retained as incident evidence for a lawful purpose (aid to a person in distress and evidence for law enforcement); and the encrypted alert carries an event class and sensor flags, not speech content. A deployed system would still owe DPDP-compliant notice (signage at instrumented locations), a retention schedule for clips and mission video, and a designated data fiduciary. These are deployment-phase obligations, noted here so that the prototype’s architecture does not have to change to meet them.

7.2 Alert integrity: why spoofing is closed

An unauthenticated alert channel would make the system an attack tool: anyone with a Rs. 500 radio could launch a police-alerting drone at will. Every LoRa packet is therefore sealed (AES-128-CTR confidentiality, truncated HMAC-SHA256 authenticity, per-node derived keys, monotonic replay counter; see §4.3); the hub drops unknown node ids, bad MACs, and stale counters before any processing. Above the radio layer, the hub reaches the drone stack through the same token-authenticated, geofence-validated API as any operator. The residual risks are stated: physical node capture (mitigated by per-node keys: one captured node compromises one pole, revocable in the registry), and radio jamming (a denial, not a spoof; multi-node corroboration and hub-side node-liveness monitoring are the future answers).

7.3 Flight law (India, Drone Rules 2021)

Flight operations are governed by **The Drone Rules, 2021**, notified by the Ministry of Civil Aviation as **G.S.R. 589(E)** on 25 August 2021 [19]. Three of the Rules' mechanisms bear directly on this project. First, **registration**: every drone must be registered on the DigitalSky platform and issued a Unique Identification Number (UIN); the prototype airframe is registered accordingly. Second, the **airspace zoning map**: DigitalSky publishes an interactive map dividing Indian airspace into **green zones** (operations up to 120 m AGL without prior permission), **yellow zones** (controlled airspace, prior permission required), and **red zones** (operations generally prohibited); all prototype flying is planned inside a green zone, and the 120 m altitude bound of the green-zone regime is enforced software-side at the API edge, with the geofence enforced both at the edge and in flight. Third, **operating category**: all prototype flying is **VLOS-only**, in an open private field, with an RC transmitter in a safety pilot's hand as the override authority (§5.7). Autonomous beyond-visual-line-of-sight response, which a deployed VanniKawachh would ultimately perform, is described in this thesis strictly as a supervised pilot-program pathway requiring regulatory engagement under the Rules, not as something the prototype does. The kit-drop payload is a first-aid kit released from ≤ 3 m over the incident point, with the fail→RTL rule of §5.5; nothing is ever dropped from transit altitude.

7.4 Dispatch restraint

The system's authority to launch an aircraft is gated three times: Stage-1 detection (recall-tuned, on the node), Stage-2 verification plus fused severity against explicit thresholds (precision-tuned, on the hub, pinned by tests asserting no dispatch below threshold), and the flight stack's own edge validation and queue admission. Every held-back incident

is still logged and visible on the dashboard, so holding back the drone never hides an incident from the police.

7.5 Radio-spectrum compliance

The LoRa alert link must itself be lawful. In India, spectrum use is administered by the **Wireless Planning and Coordination (WPC) Wing** of the Department of Telecommunications, which has **delicensed the 865 to 867 MHz band** for low-power wireless devices [21]. Indian LoRaWAN deployments ordinarily operate in this band without an operator licence, subject to the notified power and bandwidth limits. The SX1278 modules used in the prototype are 433 MHz parts; 433 MHz operation in India falls under the WPC’s separate **low-power short-range device (SRD) provisions**, which permit only very low radiated power. The prototype’s bench and field-test configuration therefore respects the applicable low-power limits, and the deployment plan standardizes on **865 to 867 MHz** hardware (SX1276-class radios) for any at-scale installation, where the delicensed band offers both legal headroom and better link budget at permitted power. This is a procurement change, not a design change: the packet format, cryptography, and gateway software of Chapter 4 are frequency-agnostic.

Chapter 8

Conclusions and Future Work

8.1 Conclusions

VanniKawachh shows that the components of an infrastructure-borne women-safety chain integrate into a single system: TinyML screening on a solar pole budget, pretrained deep audio verification on a locality hub, sealed alerting over an operator-free radio, and a safety-interlocked autonomous first response. The architecture can be proven end-to-end in simulation before any hardware is soldered. The project’s governing discipline, inherited from the v1 flight stack and extended to every new layer, is that *a claim is not a confirmation*: a Stage-1 hit is not an incident until Stage 2 and fusion say so; a packet is not an alert until its MAC and counter say so; a mode command is not a mode until the autopilot’s telemetry says so; a “finished” mission is not finished until the vehicle is disarmed on the ground. Under that discipline the full-chain rehearsal passed on the first architecture (fused severity 0.88, kit release at 3.1 m, 328 s mission, 8/8 acceptance checks, 68/68 unit cases), and the remaining work is measurement and hardening, not redesign.

Equally deliberate is what this thesis does not claim: no acoustic accuracy has been asserted that was not measured, and every number produced by a stand-in is labelled. The literature’s 97.5%-accurate wearables fail when absent; a system that aspires to replace them must not begin life with numbers that fail when examined.

8.2 Future work

1. **Phase 1 and 2 measurement campaigns:** train and flash the Stage-1 TFLM model; measure outdoor detection distance vs. SNR, Stage-1 and Stage-2 latency, end-to-end false-positive rate on street noise; measure LoRa range and loss vs. spreading factor; tune the fusion weights against bench data.
2. **Live police streaming:** RTSP/WebRTC video from the drone to the dashboard

over LTE, extending the evidence recorder of §5.6 into a real-time feed for responding officers.

3. **OpenCV victim tracking:** on-drone detection and framing of the person in distress during hover, keeping the camera and the aircraft usefully positioned without operator input.
4. **TDOA multi-node localization:** with synchronized clocks, time-difference-of-arrival across neighbouring nodes upgrades “which pole heard it” to a position estimate between poles, and provides the multi-node corroboration path sketched in §3.2.
5. **City-scale mesh:** many nodes per hub, many hubs per city; hub-to-hub LoRa or wired backhaul; fleet dispatch with per-airframe interlocks (the v1 multi-vehicle design carries over).
6. **Flight-stack modernization** (inherited from v1): pymavlink [23] / MAVSDK migration, ArduPilot 4.x SITL validation, fault-injection flight campaigns, and the Phase-3/4 hardware progression flown on the documented build.

Bibliography

- [1] R. Pavithra, R. Ahalya, and C. Shuruthika, “Discreet AI-powered women’s safety app with audio recognition and emergency automation,” in *Proc. 2026 IEEE Students Conf. Eng. Syst. (SCEECS)*, Bhopal, India, 2026, doi: 10.1109/SCEECS68810.2026.11430037.
- [2] Y. Kim, D. Jang, and S.-P. Lee, “Robust scream detection and scream temporal interval prediction using CNN-Transformer and windowing CNN,” *IEEE Access*, vol. 13, pp. 71374-71387, 2025, doi: 10.1109/ACCESS.2025.3556729.
- [3] P. Sharma and T. J. Jebaseeli, “Smart scream and panic detection system for women using AI,” in *Proc. 2025 Int. Conf. Autom. Comput. Renew. Syst. (ICACRS)*, 2025, pp. 1554-1559, doi: 10.1109/ICACRS67045.2025.11324333.
- [4] A. Srimathi, S. Jothi, and M. P. Sindhiya, “Human scream detection and analysis for crime reduction,” in *Proc. 2025 Int. Conf. Multidiscip. Sci. Comput. Intell. (ICM-SCI)*, Erode, India, 2025, pp. 1528-1534, doi: 10.1109/ICMSCI62561.2025.10894058.
- [5] P. S. Bharathi et al., “A novel IoT enabled women safety system design with emergency alert mechanism,” in *Proc. 2025 Int. Conf. Future Technol. Syst. (ICFTS)*, 2025, pp. 1-7, doi: 10.1109/ICFTS62006.2025.11031774.
- [6] R. Snehith, S. Saranya, and N. R. Reddy, “A smart device for women safety using IoT,” in *Proc. 2025 Int. Conf. Intell. Syst. Sci. (ICISS)*, 2025, pp. 277-282, doi: 10.1109/ICISS63372.2025.11076270.
- [7] Y. Kim, D. Jang, and J. Lee, “Development of scream detection system with large-scale scream dataset,” in *Proc. 2024 Int. Conf. Inf. Commun. Technol. Converg. (ICTC)*, Jeju, South Korea, 2024, pp. 590-593, doi: 10.1109/ICTC62082.2024.10826757.
- [8] A. A. Fime, M. Ashikuzzaman, and A. Aziz, “Audio signal-based danger detection using signal processing and deep learning,” *Expert Systems with Applications*, vol. 237, art. no. 121646, Mar. 2024, doi: 10.1016/j.eswa.2023.121646.

- [9] R. V. Hadkar et al., “An efficient IoT-enabled women safety device,” in *Proc. 2024 Int. Conf. Autom. Comput. Renew. Syst. (ICACRS)*, 2024, pp. 425-431, doi: 10.1109/ICACRS62842.2024.10841797.
- [10] G. Uganya et al., “Smart women safety device using IoT and GPS tracker,” in *Proc. 2023 Int. Conf. Comput. Electron. Biomed. Syst. (ICCEBS)*, Chennai, India, 2023, pp. 1-6, doi: 10.1109/ICCEBS58601.2023.10449302.
- [11] S. Potturi et al., “An SOS women safety device: Automatic emergency alerts and geo-location sharing with GSM/GPS integration,” in *Proc. 2026 Int. Conf. Intell. Comput. Netw. Syst. (IC-ICNS)*, Bhubaneswar, India, 2026, pp. 1-6, doi: 10.1109/IC-ICNS68863.2026.11537940.
- [12] G. Ciaburro and V. Puyana-Romero, “Sound event detection in smart cities: A systematic review of methods, datasets, and applications,” *Big Data and Cognitive Computing*, vol. 10, no. 3, art. no. 83, 2026, doi: 10.3390/bdcc10030083.
- [13] Q. Kong, Y. Cao, T. Iqbal, Y. Wang, W. Wang, and M. D. Plumbley, “PANNs: Large-scale pretrained audio neural networks for audio pattern recognition,” *IEEE/ACM Trans. Audio, Speech, Lang. Process.*, vol. 28, pp. 2880-2894, 2020, doi: 10.1109/TASLP.2020.3030497; `panns-inference` package.
- [14] R. David et al., “TensorFlow Lite Micro: Embedded machine learning for TinyML systems,” in *Proc. Mach. Learn. Syst. (MLSys)*, vol. 3, 2021, pp. 800-811; documentation and the `micro_speech` example, <https://www.tensorflow.org/lite/microcontrollers>.
- [15] ArduPilot Development Team, “ArduPilot documentation” (firmware, SITL, and failsafe documentation), <https://ardupilot.org>; and MAVLink Development Team, “MAVLink developer guide” (HEART-BEAT, SET_MODE, COMMAND_LONG/MAV_CMD_DO_SET_MODE, MAV_CMD_DO_SET_SERVO), <https://mavlink.io/en/>. Accessed Jul. 2026.
- [16] LoRa Alliance, “LoRaWAN L2 1.0.4 specification,” LoRa Alliance Technical Committee, 2020. [Online]. Available: https://lora-alliance.org/resource_hub/lorawan-104-specification-package/.
- [17] National Institute of Standards and Technology, “Advanced Encryption Standard (AES),” FIPS PUB 197, Nov. 2001, doi: 10.6028/NIST.FIPS.197.
- [18] H. Krawczyk, M. Bellare, and R. Canetti, “HMAC: Keyed-hashing for message authentication,” IETF RFC 2104, Feb. 1997.

- [19] Ministry of Civil Aviation, Government of India, “The Drone Rules, 2021,” notification G.S.R. 589(E), *Gazette of India*, 25 Aug. 2021; DigitalSky platform (UIN registration and airspace zone map). <https://digitalsky.dgca.gov.in>. Accessed 2026-07-06.
- [20] Government of India, “The Digital Personal Data Protection Act, 2023,” Act No. 22 of 2023, *Gazette of India*, 11 Aug. 2023.
- [21] Wireless Planning and Coordination Wing, Department of Telecommunications, Government of India: delicensing of the 865-867 MHz band for low-power wireless devices (G.S.R. 1048(E), 2005) and low-power short-range-device (SRD) provisions applicable to 433 MHz operation.
- [22] DroneKit-Python 2.9.2. <https://github.com/dronekit/dronekit-python>. Accessed 2026-07-06.
- [23] pymavlink. <https://github.com/ArduPilot/pymavlink>. Accessed 2026-07-06.
- [24] dronekit-sitl 3.3.0. <https://github.com/dronekit/dronekit-sitl>. Accessed 2026-07-06.
- [25] FastAPI. <https://fastapi.tiangolo.com>. Accessed 2026-07-06.
- [26] React 18. <https://react.dev>. Accessed 2026-07-06.
- [27] Leaflet 1.9. <https://leafletjs.com>. Accessed 2026-07-06.
- [28] OpenStreetMap. <https://www.openstreetmap.org>. Accessed 2026-07-06.
- [29] Pixhawk hardware reference. <https://pixhawk.org>. Accessed 2026-07-06.
- [30] Raspberry Pi 5 product documentation. <https://www.raspberrypi.com/documentation/>. Accessed 2026-07-06.
- [31] Espressif ESP32-S3 technical reference manual. <https://www.espressif.com>. Accessed 2026-07-06.
- [32] Semtech SX1276/77/78/79 LoRa transceiver datasheet. <https://www.semtech.com>. Accessed 2026-07-06.
- [33] InvenSense/TDK INMP441 omnidirectional I2S MEMS microphone datasheet. <https://invensense.tdk.com>. Accessed 2026-07-06.
- [34] U.S. FAA, Part 107 (Small Unmanned Aircraft Systems). https://www.faa.gov/uas/commercial_operators. Accessed 2026-07-06.

Appendix A

LoRa Alert Packet Wire Format

25 bytes total, which fits in one LoRa frame (Table A.1).

Table A.1: Sealed alert packet wire format (25 bytes).

Offset	Size	Field
0	2	magic "VK"
2	1	version (1)
3	2	node_id (uint16 BE), cleartext (selects the key)
5	4	counter (uint32 BE), cleartext (CTR nonce + replay)
9	8	AES-128-CTR ciphertext of payload
17	8	MAC: HMAC-SHA256(node_key, header + ciphertext)[:8]

Payload (before encryption): event uint8 (1 = scream, 2 = help_keyword, 3 = cry, 4 = crash); confidence uint8 (Stage-1, 0 to 255); PIR uint8 (0/1); light uint8 (LDR 0 to 255, 0 = dark); battery uint8 (%); 3 reserved bytes.

Per-node key = HMAC-SHA256(master_key, "node:<id>")[16]. Rejection rules at the hub: bad length, bad magic/version, bad MAC, unknown node_id, counter $\leq$ last accepted.

Appendix B

Configuration Reference

Table B.1: Hub configuration (`hub/config.py`).

Variable	Default	Governs
HUB_MASTER_KEY	dev key (must be set in deployment)	AES-128 master key (hex)
VERIFY_THRESHOLD	0.50	Min Stage-2 audio score to count as distress
DISPATCH_THRESHOLD	0.60	Min fused severity to launch the drone
CLIP_WAIT_S	8.0	Wait for the WiFi clip before degraded scoring
GATEWAY_PORT	/ COM3 / 115200	Gateway ESP32 serial
GATEWAY_BAUD		
DRONE_API_URL	/ localhost:8000 / un-	Drone stack endpoint
DRONE_API_TOKEN	set	
NODES_FILE / CLIPS_DIR	hub/nodes.json / hub/clips	Registry and clip storage
CLIP_SERVER_PORT	8990	Node clip upload server

Table B.2: Flight-stack configuration (`flight_core/config.py`, v2 additions).

Variable	Default	Governs
PAYLOAD_SERVO_CHANNEL	9 (Pixhawk AUX OUT 1)	Kit-release servo output
PAYLOAD_OPEN_PWM	/ 1900 / 1100	Hook open / closed
PAYLOAD_HOLD_PWM		
PAYLOAD_DROP_ALT	3.0 m	Descend-to altitude before release

The full v1 flight-stack configuration table (connection string, home, geofence, battery thresholds, GPS debounce, stall timeout, queue bounds, API token, SITL_MODE) is unchanged and lives in `docs/SYSTEM_DOCUMENTATION.md` §7.

Appendix C

Test Inventory

Tier 1: 68 collected cases. `tests/test_units.py` (47): 10 failsafe-arbiter cases (thresholds, escalation, no-downgrade, debounce $\times 2$, fire-once, geofence, timeout, healthy baseline); 7 queue cases (priority, depth cap, prune, cancel $\times 3$, worker bookkeeping); 2 persistence cases; validation cases covering every rejection bound (coordinates, altitude, hover, priority, waypoints, including parametrized edge values); 3 configuration cases. `tests/test_hub.py` (14): packet seal/unseal round-trip; MAC tamper rejection; replay-counter rejection; registry lookup + unknown-node drop; fusion scoring and priority mapping; pipeline gating (no dispatch below verify/dispatch thresholds); degraded no-clip scoring; dispatcher payload shape. `tests/test_obstacle_avoidance.py` (7): keep-out routing geometry.

Tier 2 (`tests/test_full_mission.py`): the eight-property SITL acceptance flight of §6.1, runnable on any machine.

Tier 3 (`scripts/demo_phase0.py`): the zero-hardware full-chain rehearsal of §6.3 (sensing sim $\rightarrow$ hub $\rightarrow$ SITL flight $\rightarrow$ kit drop), exit 0 on a completed mission.